

FinFlow: Automated Credit Risk Management & ROI Optimization

Project Lead: Yue Dong (Lead Machine Learning Engineer)

Date: 2026-02-05

1. Business Objectives

In the traditional lending landscape, manual credit reviews are slow, subjective, and prone to human error. FinFlow was designed as a data-driven solution to achieve three primary strategic goals:

- **Operational Efficiency:** Automate over **70%** of routine loan applications, allowing credit officers to focus on high-complexity cases.
 - **Net Profit Maximization:** Use mathematical modeling to identify the "Sweet Spot" between interest income and default loss.
 - **Systemic Stability:** Establish a production-grade monitoring framework to detect market shifts (Data Drift) in real-time.
-

2. Variable Definitions & Feature Engineering

The model utilizes both raw financial data and engineered "ratio features" to capture the nuanced repayment capacity of small businesses.

Core Feature Set

| Category | Variable | Business Rationale |
|----------------|-------------------|---|
| Solvency | loan_to_income | Ratio of total debt to annual income (The #1 risk predictor). |
| Leverage | rev_to_loan_ratio | Measures how much revenue is generated for every \$1 borrowed. |
| Credit History | credit_tier | Binned categorical rankings (Elite, Average, Subprime). |
| Efficiency | rev_per_employee | Revenue per head; reflects management quality and overhead risk. |
| Market Segment | industry | Weights risk based on sector-specific volatility (e.g., Retail vs. Tech). |

3. Methodology

The project follows the standard **Machine Learning Life Cycle (MLOps)**, divided into five distinct phases:

1. **Data Quality Audit (EDA):** Cleaned 5,000+ historical records and addressed a 15% class imbalance.
 2. **End-to-End Pipeline:** Built a robust feature pipeline using SimpleImputer for missing values and OneHotEncoder for categorical alignment.
 3. **Model Selection:** Evaluated Logistic Regression vs. Random Forest vs. XGBoost; **Random Forest** was selected for its superior performance in capturing non-linear risk patterns.
 4. **Strategy Optimization:** Utilized SHAP for model explainability and plotted **Profit-Maximization Curves**.
 5. **Engineering Deployment:** Modularized experimental code into production-ready Python scripts with built-in validation.
-

4. Key Findings & Business Strategy

Through rigorous testing and simulation, the project yielded the following outcomes:

- **Performance:** The champion model (Random Forest) achieved an **81% Precision** rate, effectively flagging the vast majority of potential defaults.
- **The 0.29 Pivot Point:** Our research proved that the standard 0.5 probability threshold was sub-optimal. The **0.29 threshold** was identified as the point of maximum net recovery:
 - **Prob < 0.20:**  **Auto-Approve** (High velocity, low risk).
 - **0.20 <= Prob < 0.29:**  **Manual Review** (Edge cases for expert audit).
 - **Prob >= 0.29:**  **Auto-Reject** (Immediate capital protection).

5. Production Practice & MLOps

To ensure stability in a live environment, FinFlow includes a comprehensive automation and monitoring suite:

- **Weekly Pipeline (`weekly_run.py`):** Automatically processes new batches of applications every week.
- **Stability Monitoring (PSI):** Real-time tracking of the Population Stability Index to prevent model decay.

$$PSI = \sum (Actual\% - Expected\%) \times \ln \left(\frac{Actual\%}{Expected\%} \right)$$

- **Latest Result: 0.0003** (Indicates near-perfect alignment with training data).
 - **Automated Auditing:** Generates ROI reports for every run, quantifying "Potential Loss Avoided."
-

6. Future Roadmap

While the system is currently live and stable, the following steps are planned to maintain a competitive edge:

A/B Testing Strategy

- **Experiment Design:** Allocate 10% of traffic to an "Aggressive Group" (Threshold set at 0.35).
- **Goal:** Observe if the increased interest income from higher approval volume offsets the marginal increase in default costs.

System Evolution

- **Real-Time API:** Transition from weekly batch processing to an API-based "Instant Decision" architecture.
- **Alternative Data:** Integrate external datasets (e.g., utility payments or social sentiment) to further refine the credit profiles of small businesses.